

1

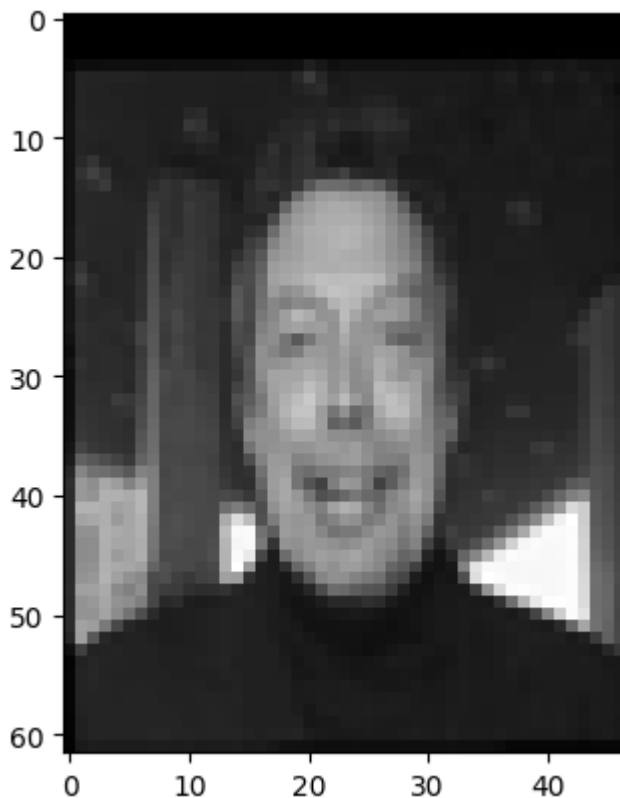
```
In [ ]: from IPython.display import Latex
```

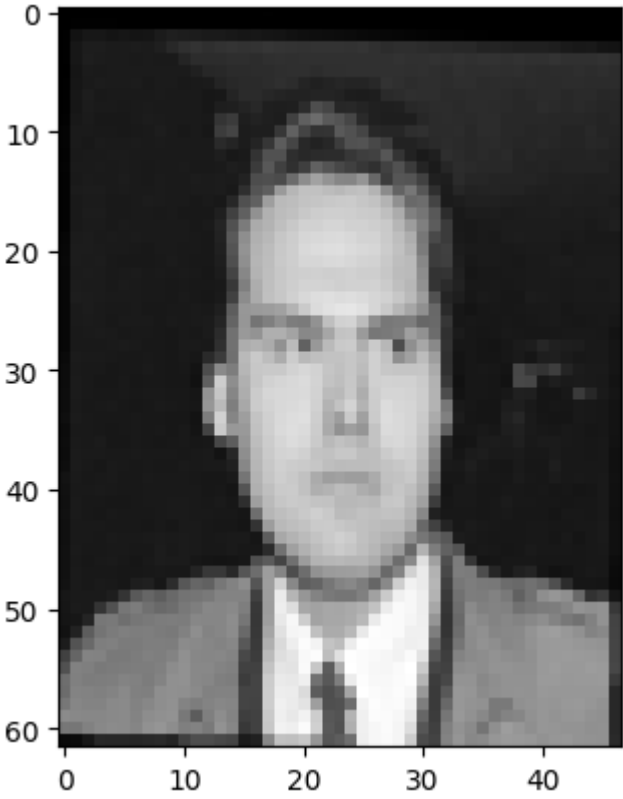
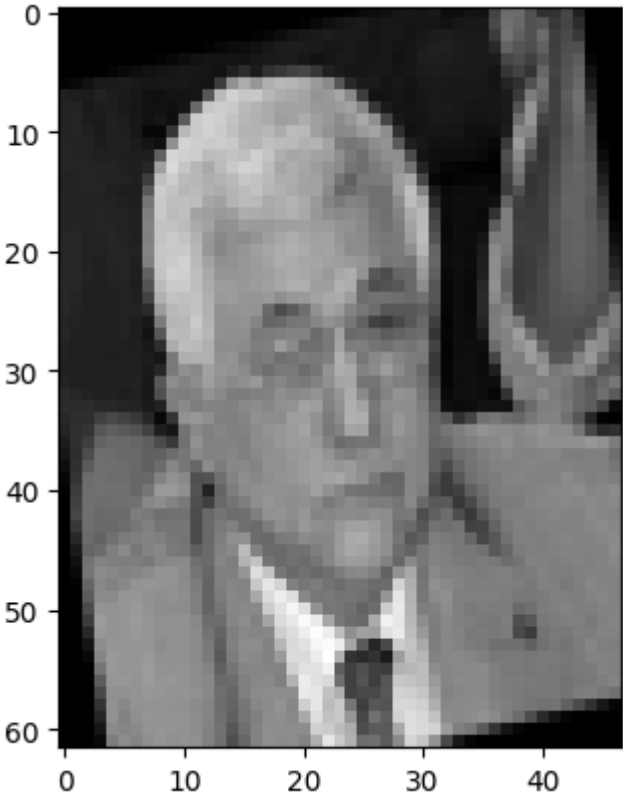
```
In [ ]: #Downloading dataset
from sklearn import datasets
dataset = datasets.fetch_lfw_people()
X = dataset['data']
#check data
print(X.data.shape)

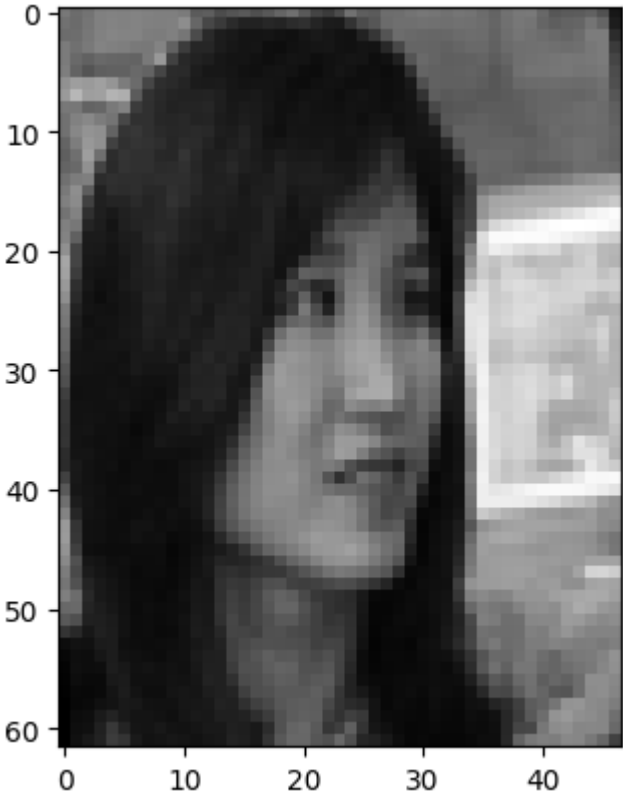
(13233, 2914)
```

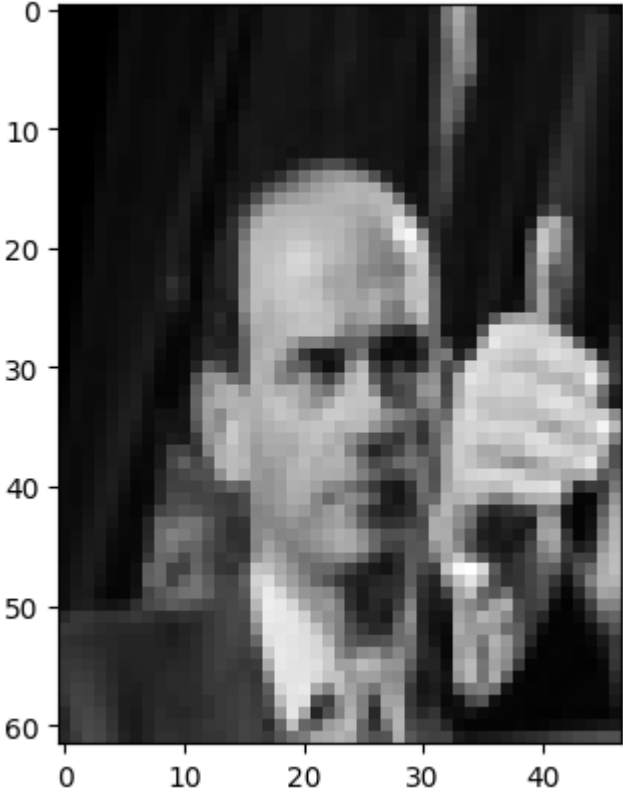
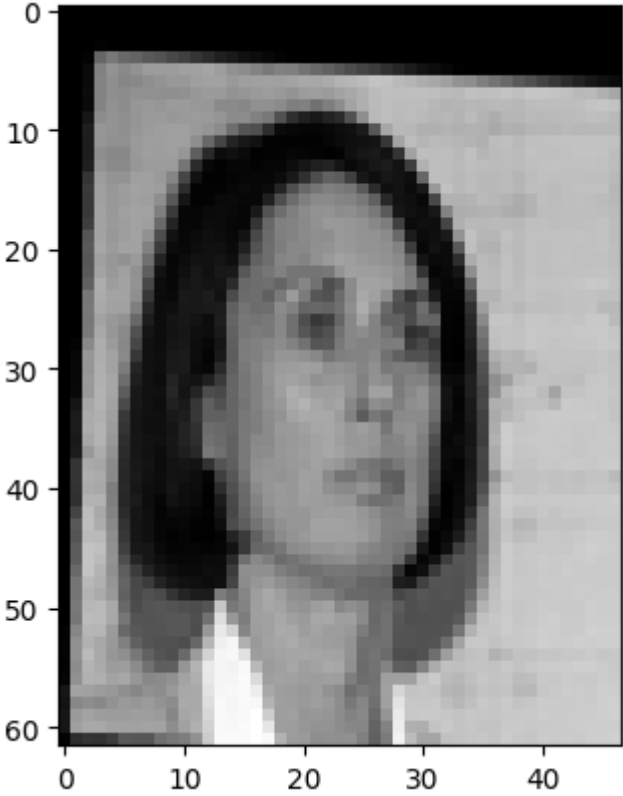
```
In [ ]: #Relevant imports
import matplotlib
import numpy as np
from matplotlib import pyplot
from numpy import linalg
```

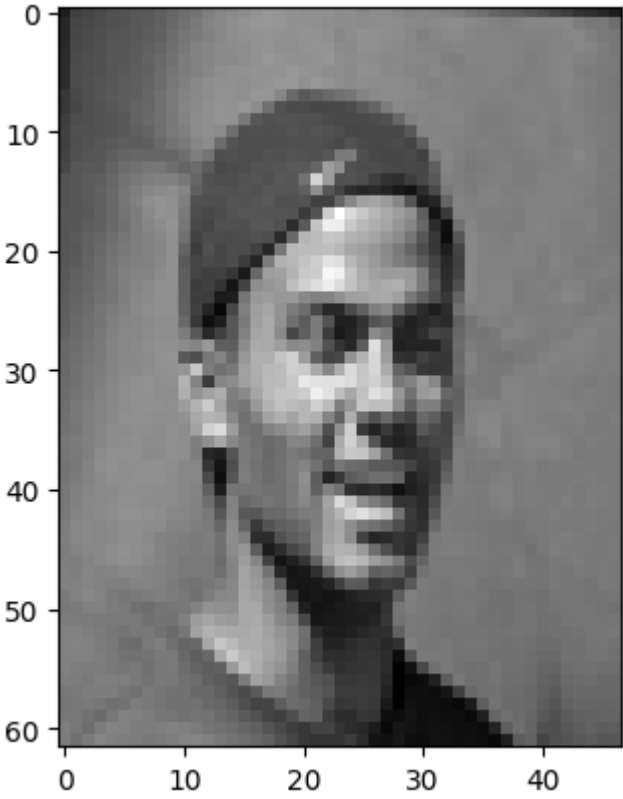
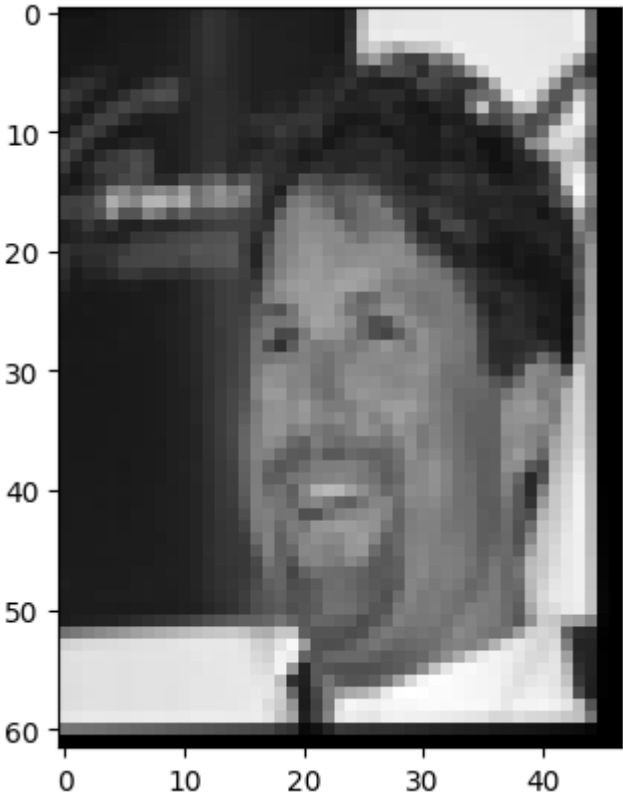
```
In [ ]: #1.a
for i in range(20):
    matplotlib.pyplot.imshow(X[i].reshape((62,47)), cmap=matplotlib.pyplot.cm.gray)
    matplotlib.pyplot.savefig(f'files/1_a={i}.png')
    matplotlib.pyplot.show()
```

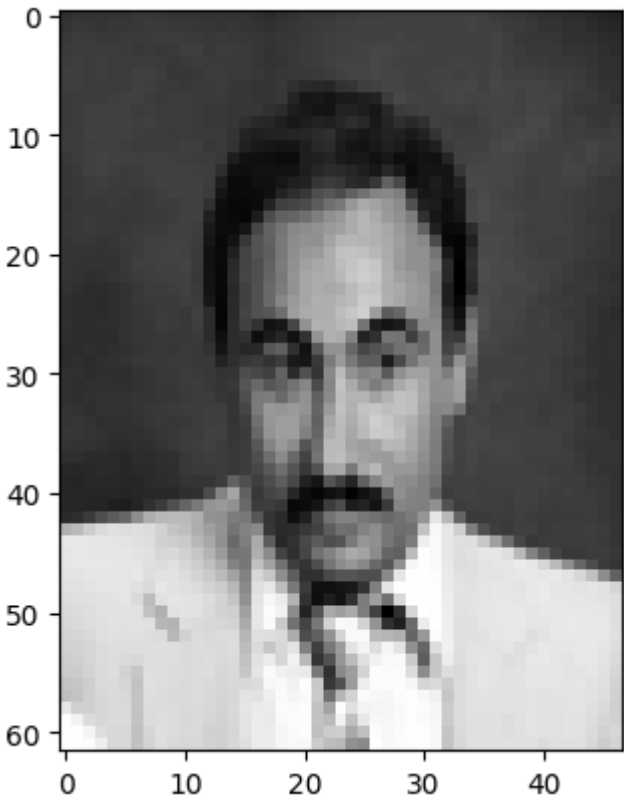






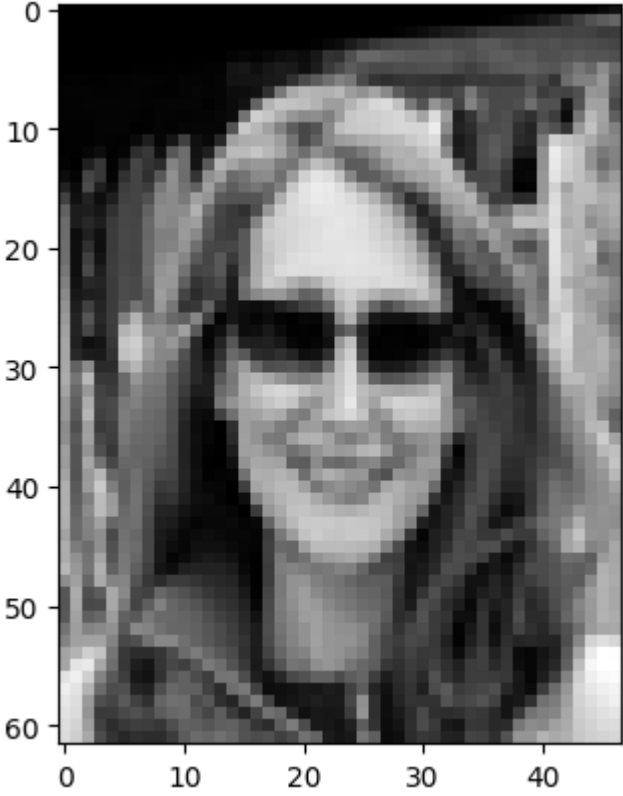








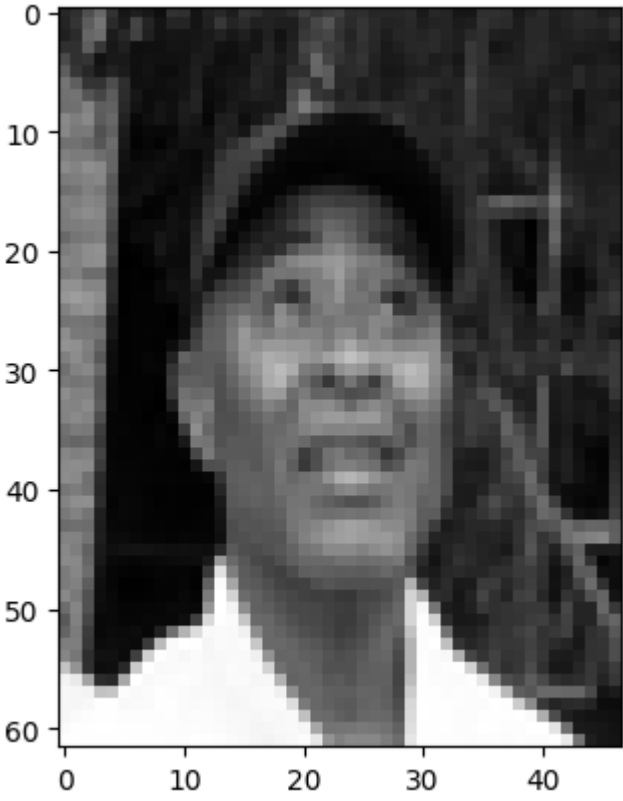
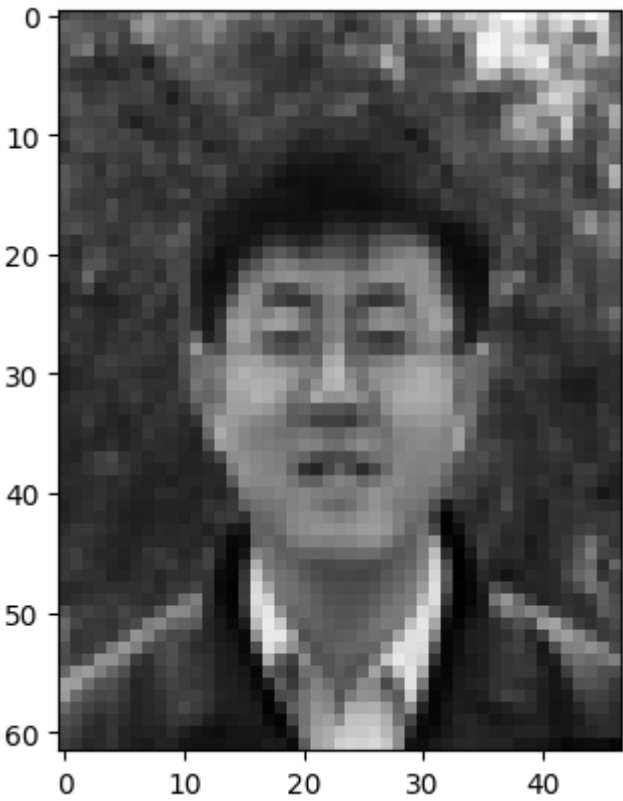
















```
In [ ]: display(Latex(r"\newpage"))
```

\newpage

```
In [ ]: #1.b
average = X.mean(axis = 0)
print(average)
for i in range(X[0].size):
    X[i]-=average
matplotlib.pyplot.imshow(average.reshape((62,47)), cmap=matplotlib.pyplot.cm.gray )
matplotlib.pyplot.savefig(f'files/1_b=average.png')
matplotlib.pyplot.show()
```

```
[0.13771963 0.15683803 0.17031944 ... 0.2221487 0.20515166 0.18089974]
```



```
In [ ]: display(Latex(r"\newpage"))
```

\newpage

```
In [ ]: #1.c
eig_vectors, eig_values, _ = np.linalg.svd(X.T@X)
```

```
In [ ]: #1.c

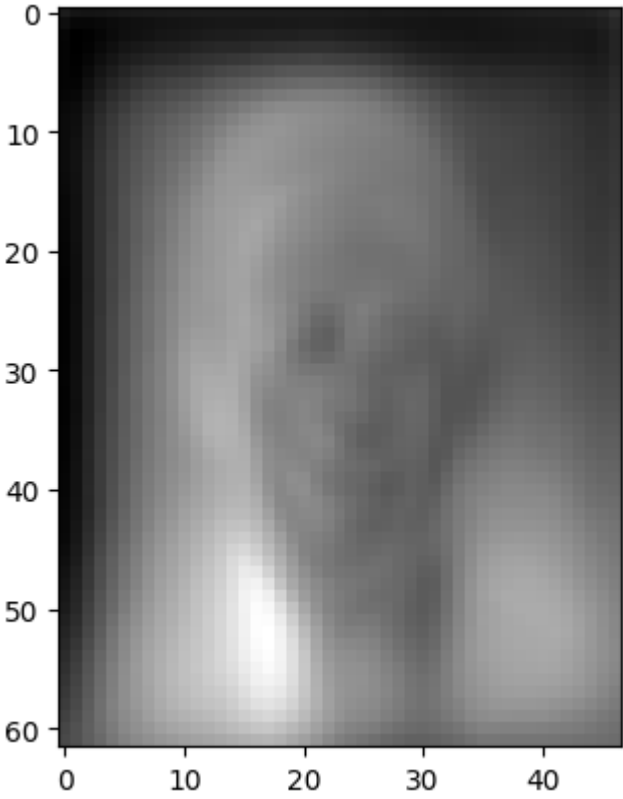
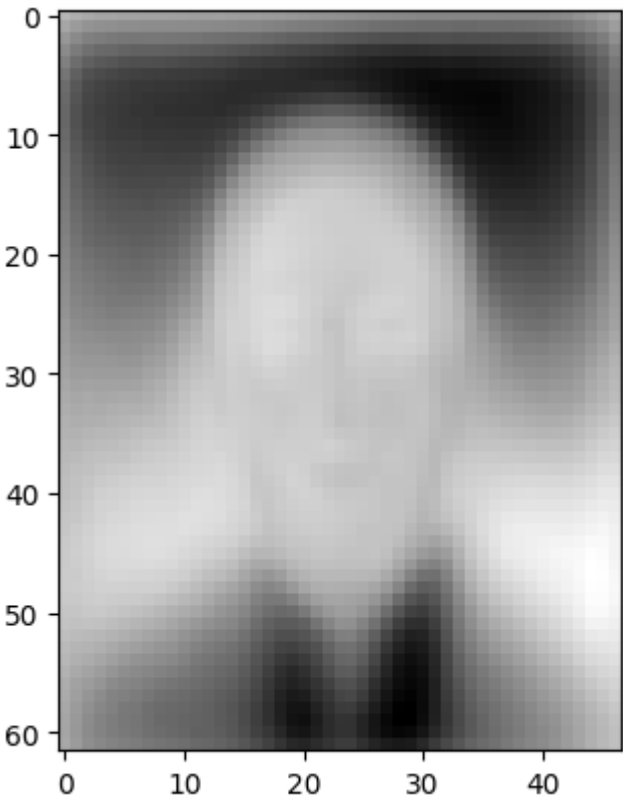
def projectTopKEigenVectors(X, k, eig_values, eig_vectors):
    idx = eig_values.argsort()[::-1]
    eig_values = eig_values[idx]
    eig_vectors = eig_vectors[:,idx]

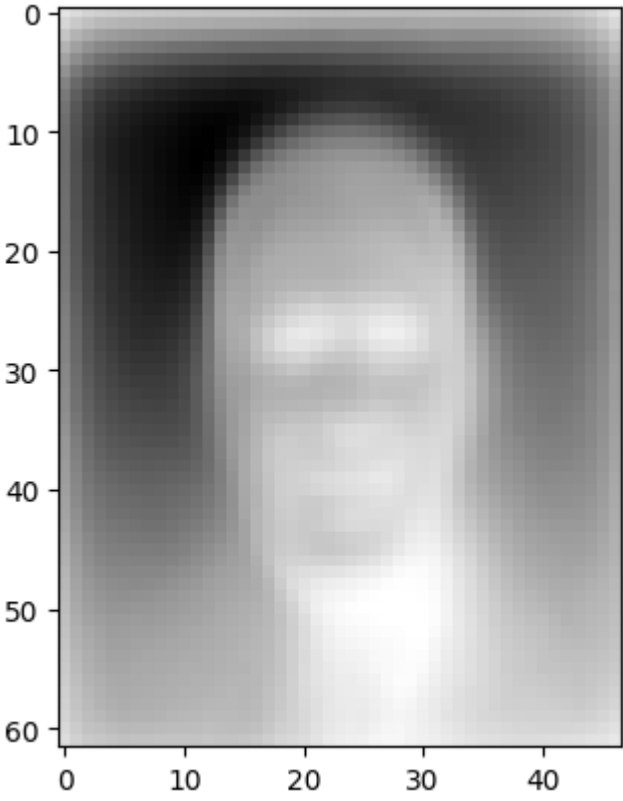
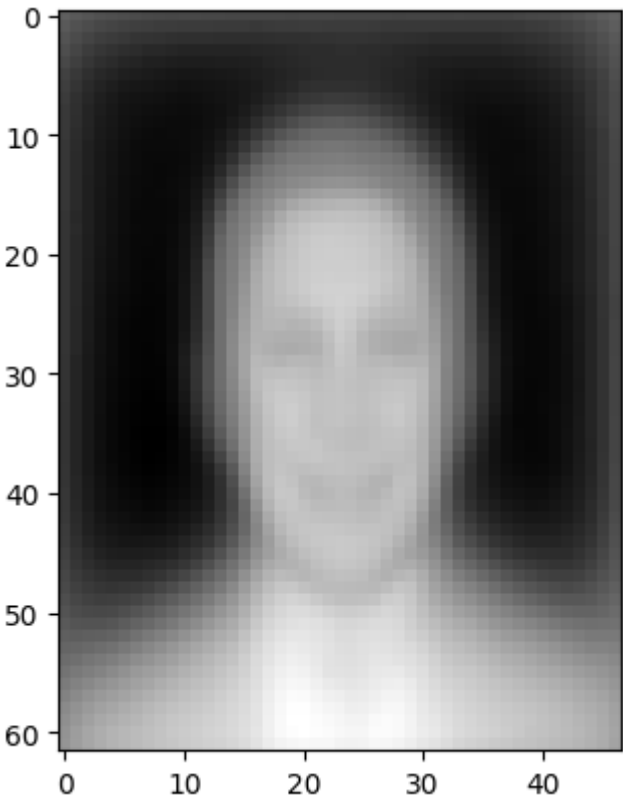
    k_eig_vectors = np.zeros(eig_vectors.shape)
    k_eig_vectors[:, :k] = eig_vectors[:, :k]



    X_k = X @ k_eig_vectors
    X_recon = X_k @ k_eig_vectors.transpose()

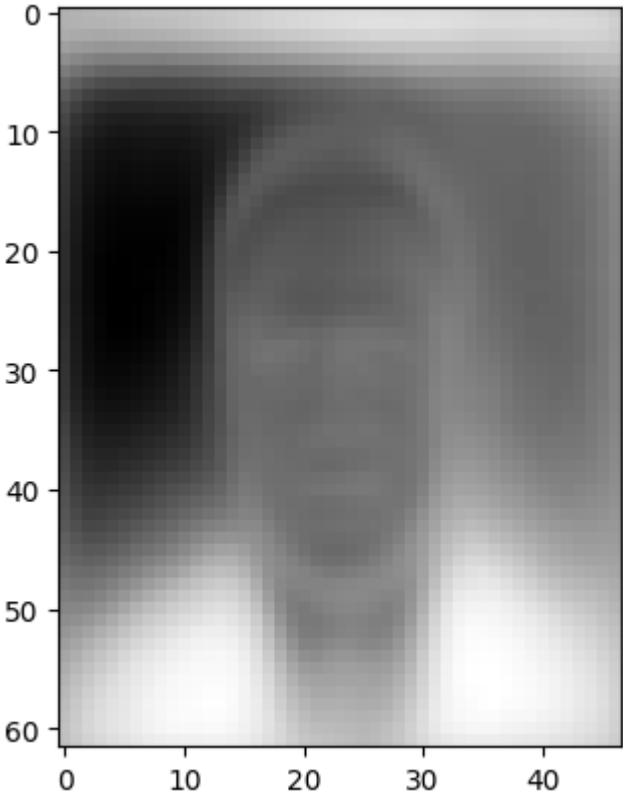
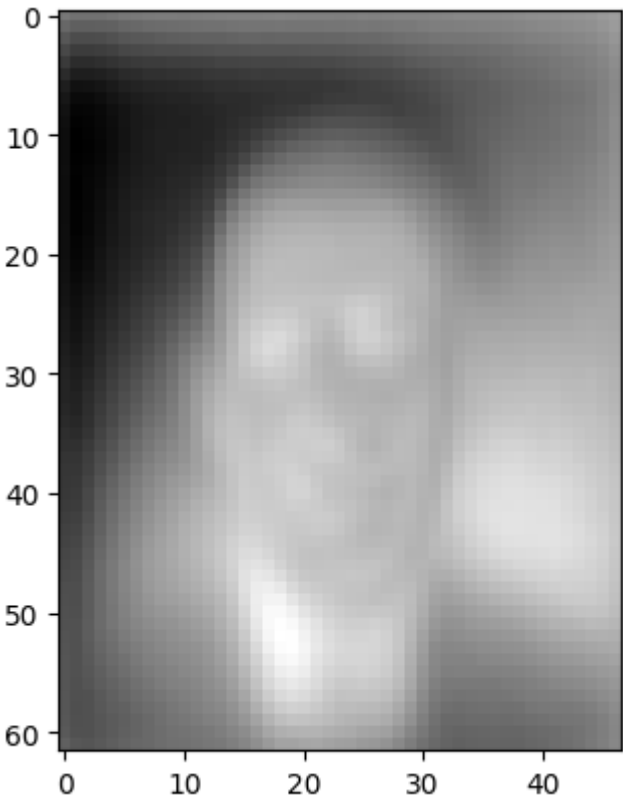
    for i in range(20):
        matplotlib.pyplot.imshow(X_recon[i].reshape((62,47)), cmap=matplotlib.pyplot.cm.gray)
        matplotlib.pyplot.savefig(f'files/1_c_top20_nr.{i+1}_k={k}.png')
        matplotlib.pyplot.show()

projectTopKEigenVectors(X, 10, eig_values, eig_vectors)
projectTopKEigenVectors(X, 100, eig_values, eig_vectors)
projectTopKEigenVectors(X, 1000, eig_values, eig_vectors)
```

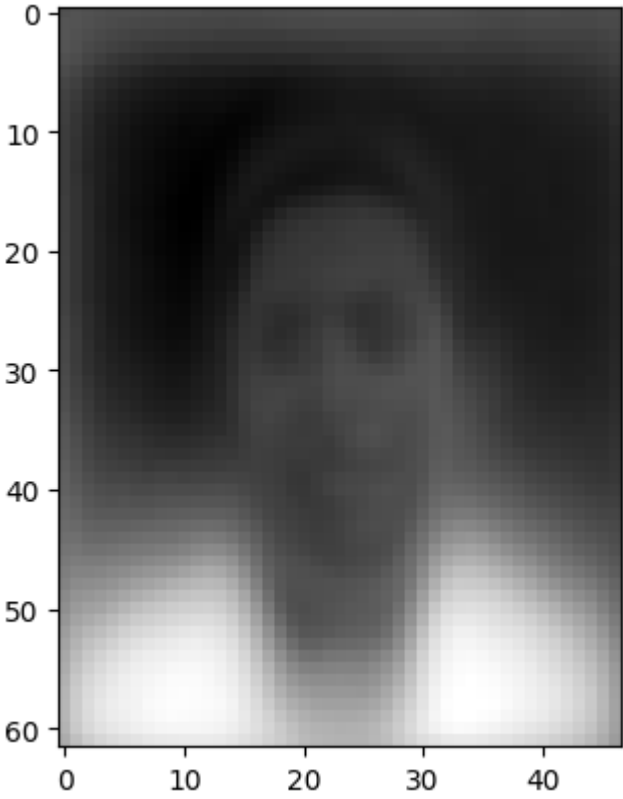
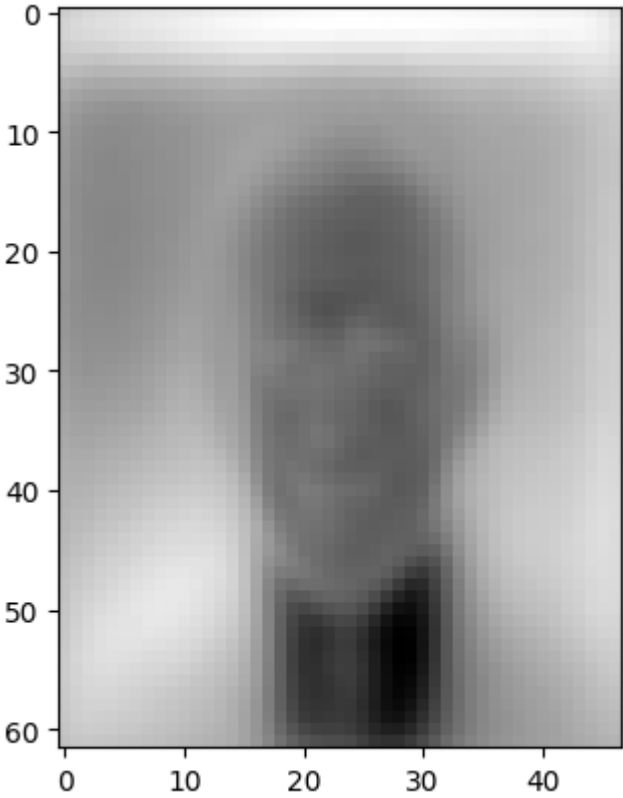






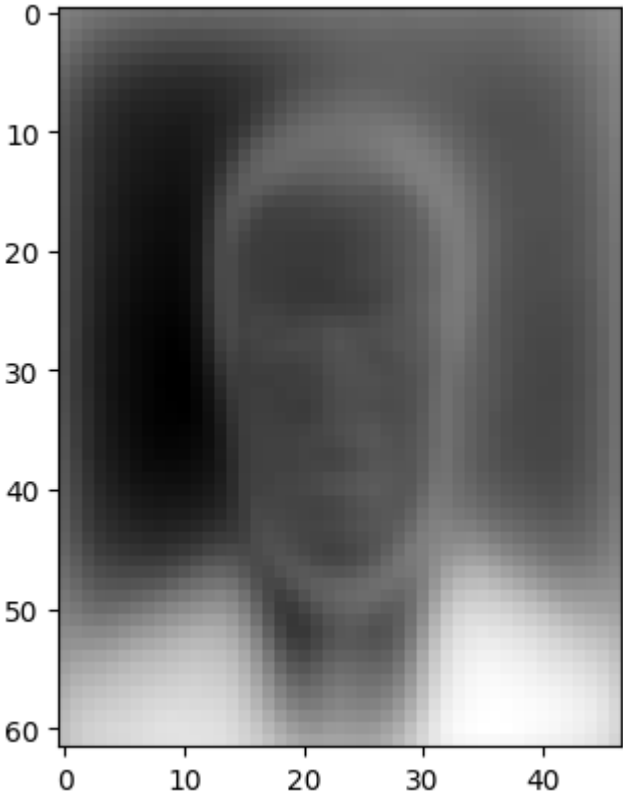



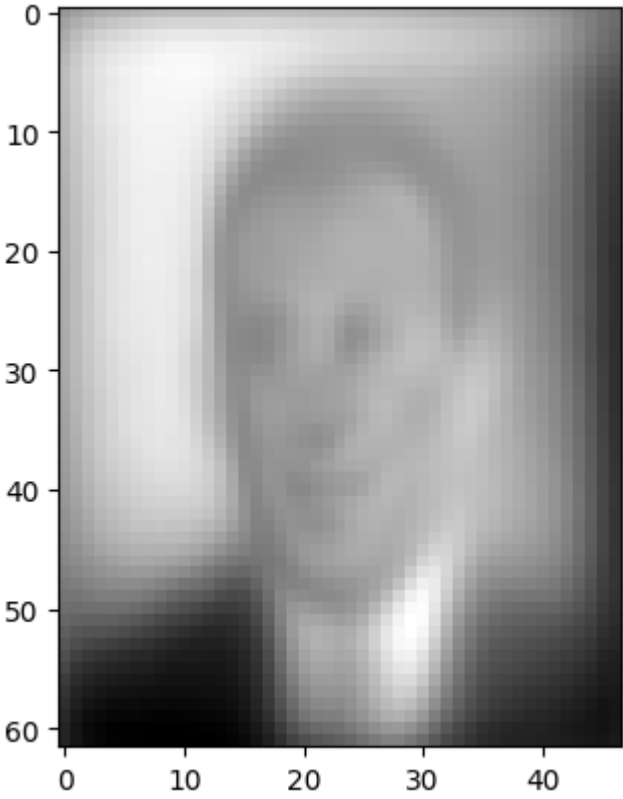



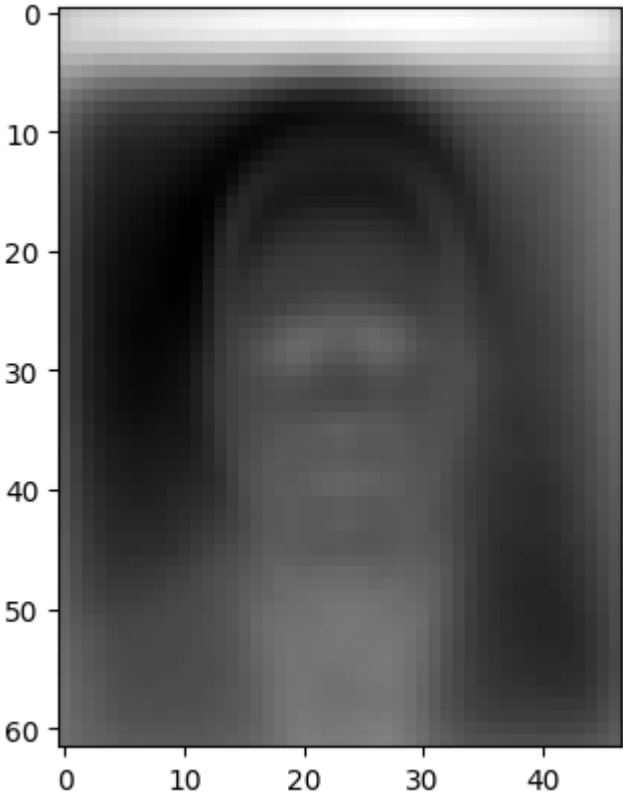
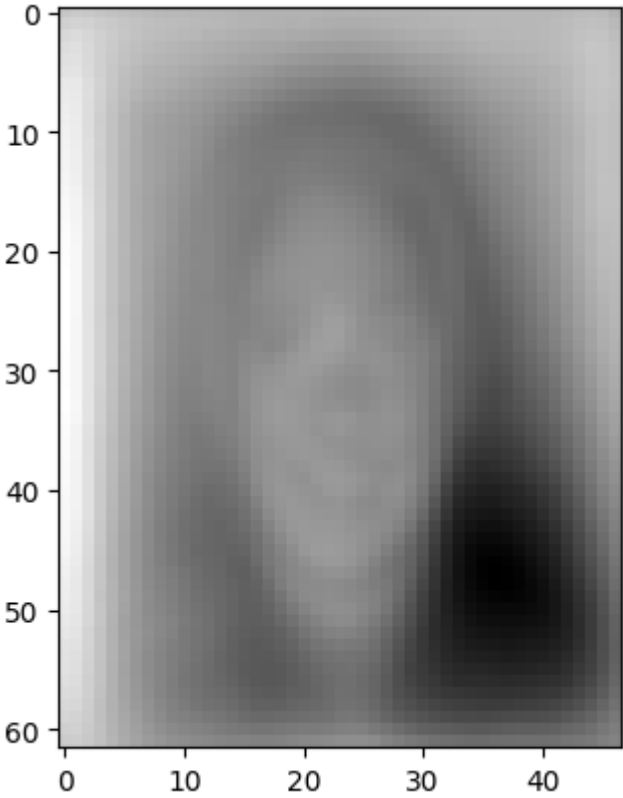


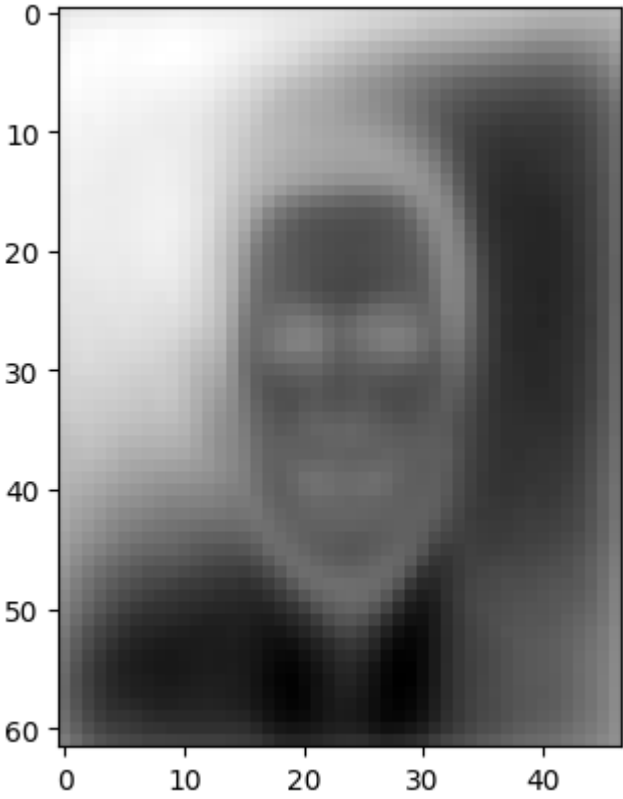
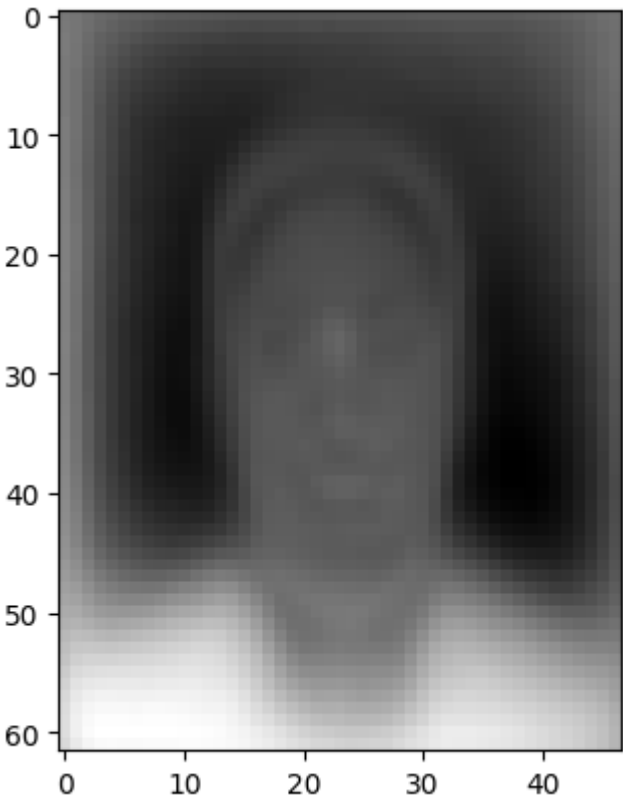



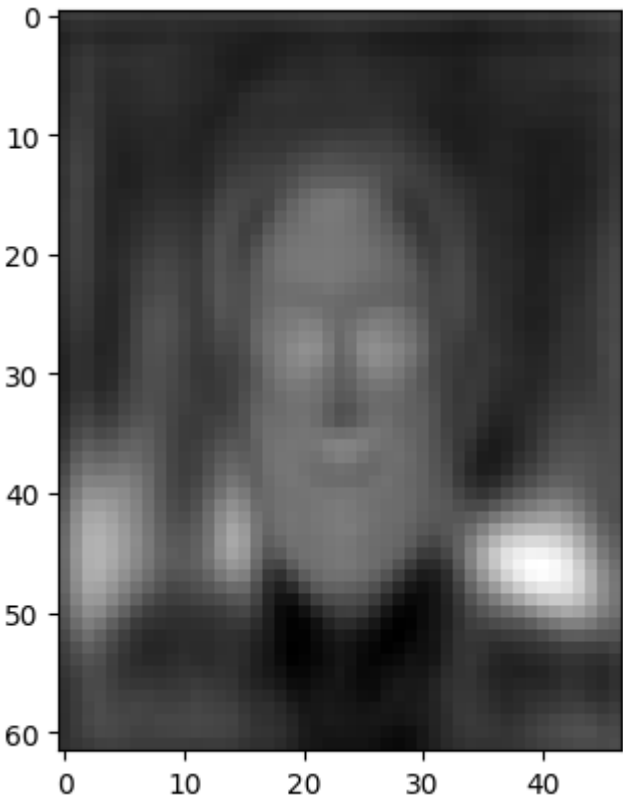



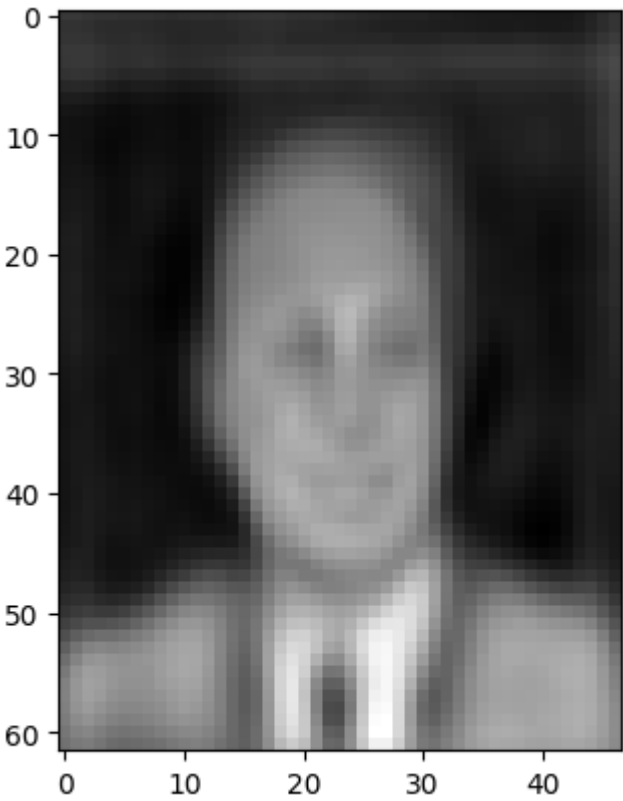




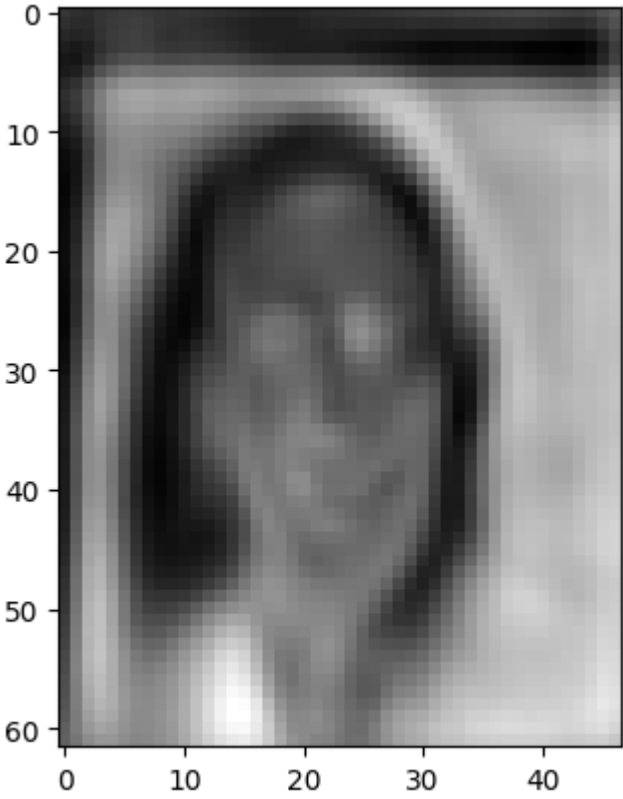
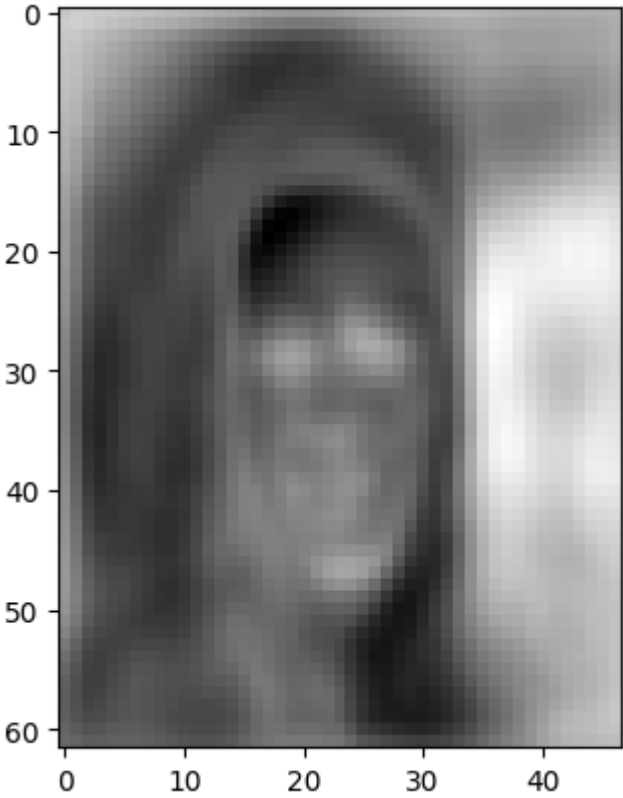






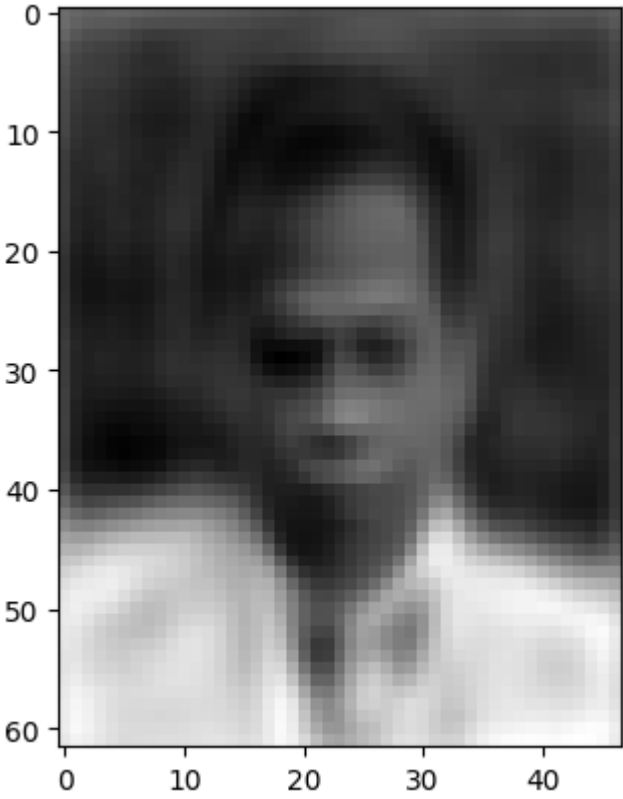











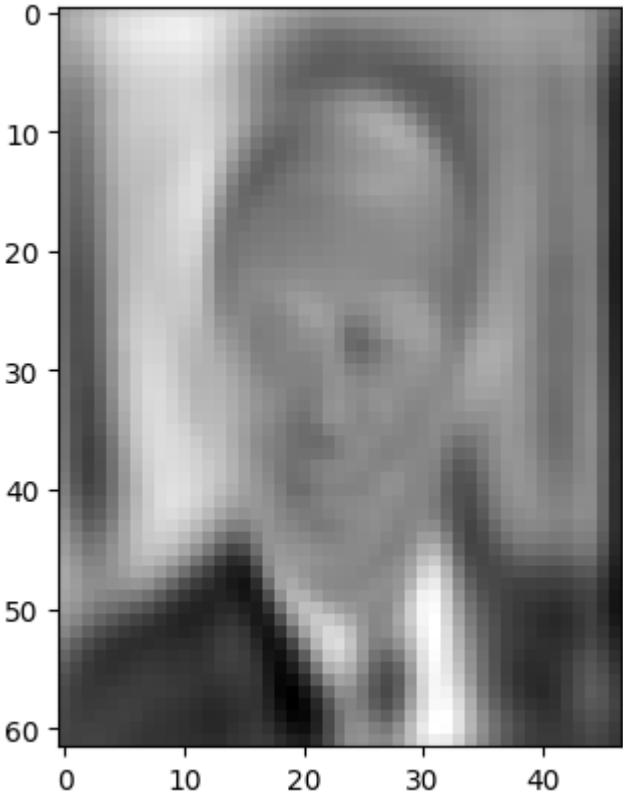



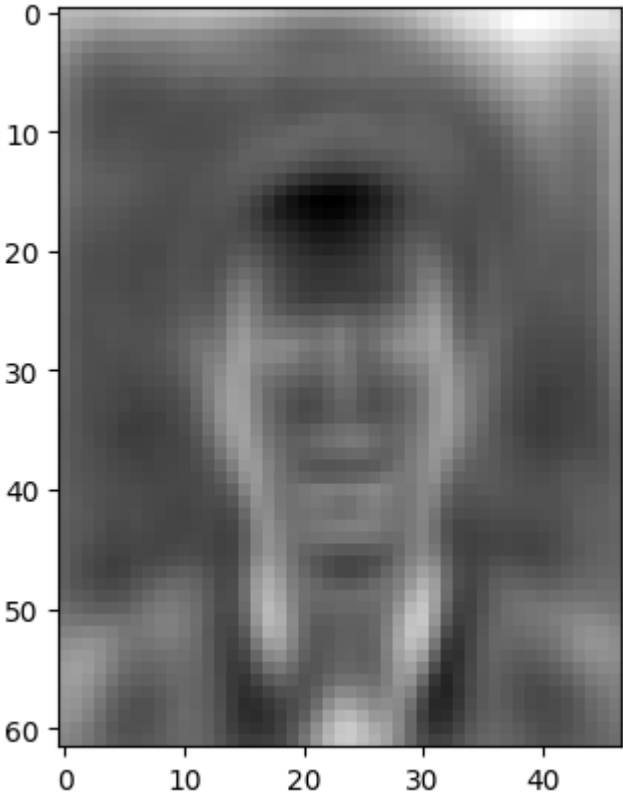









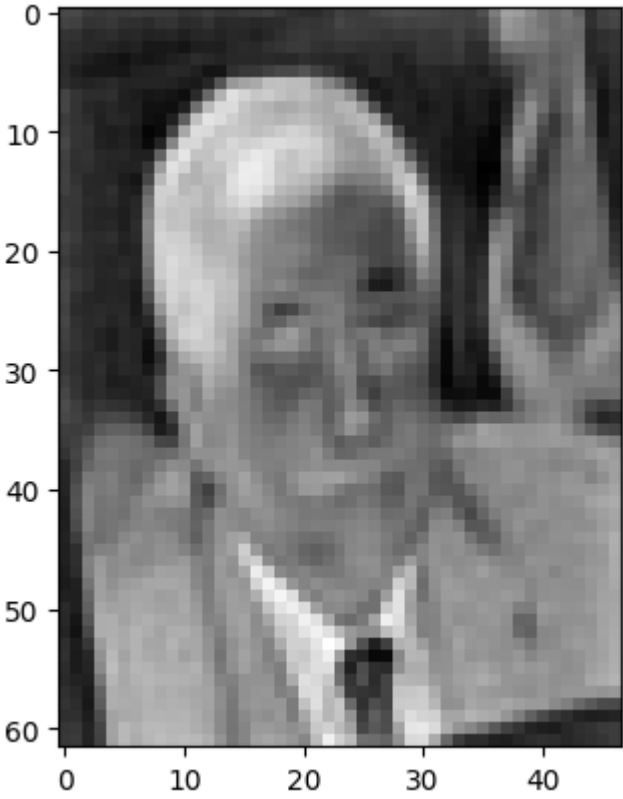
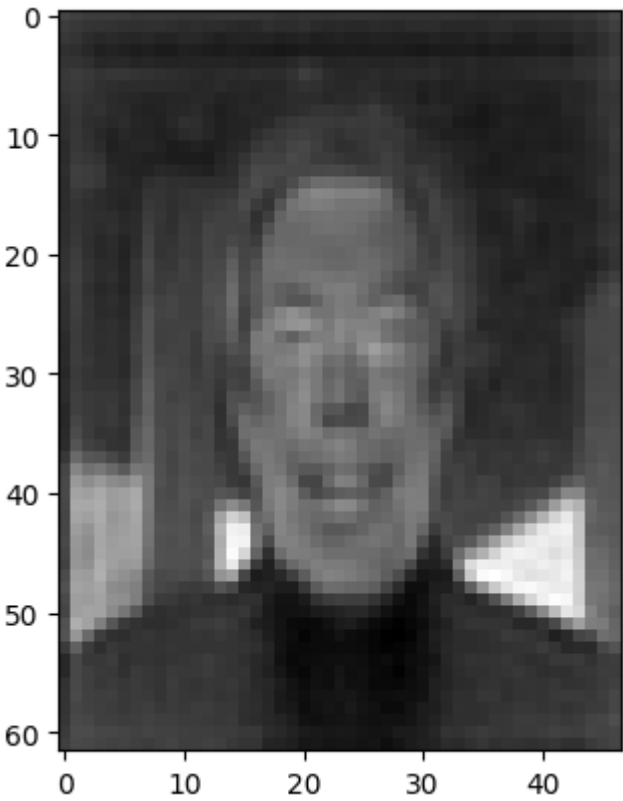



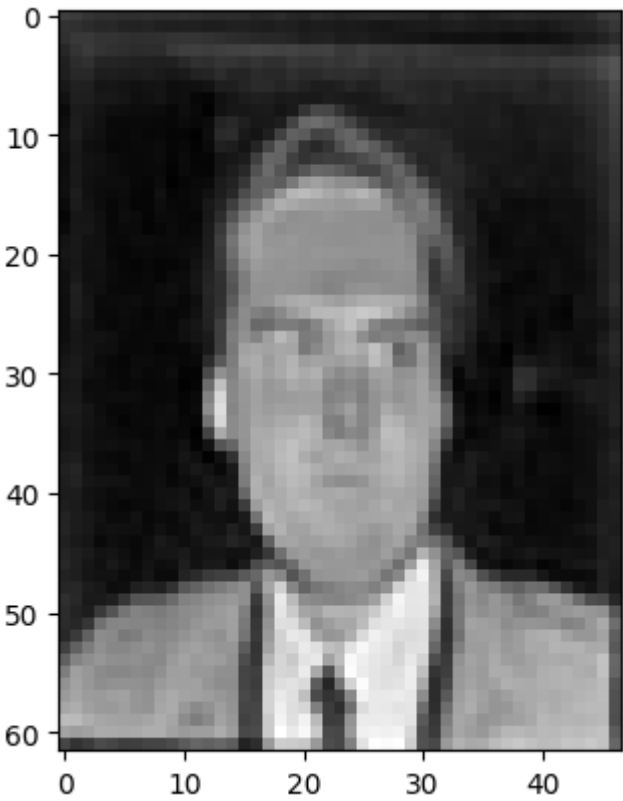






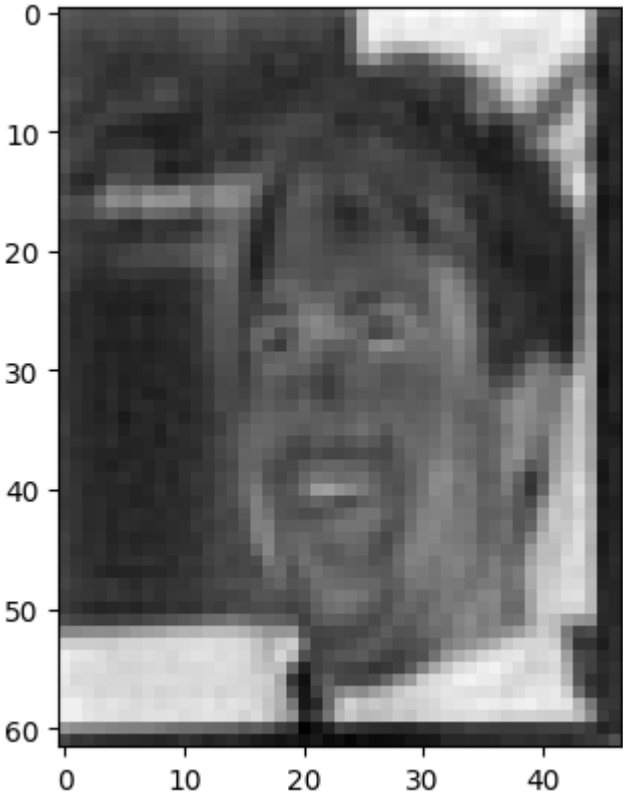
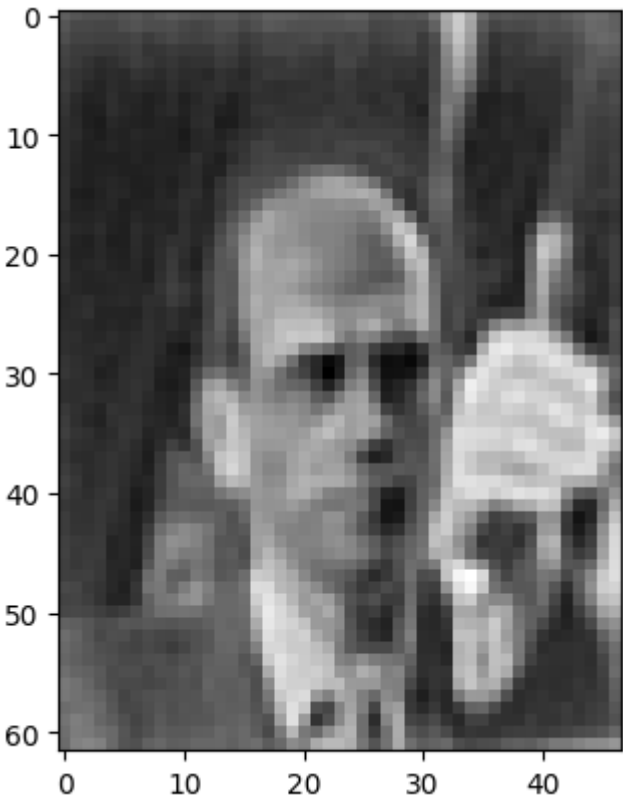






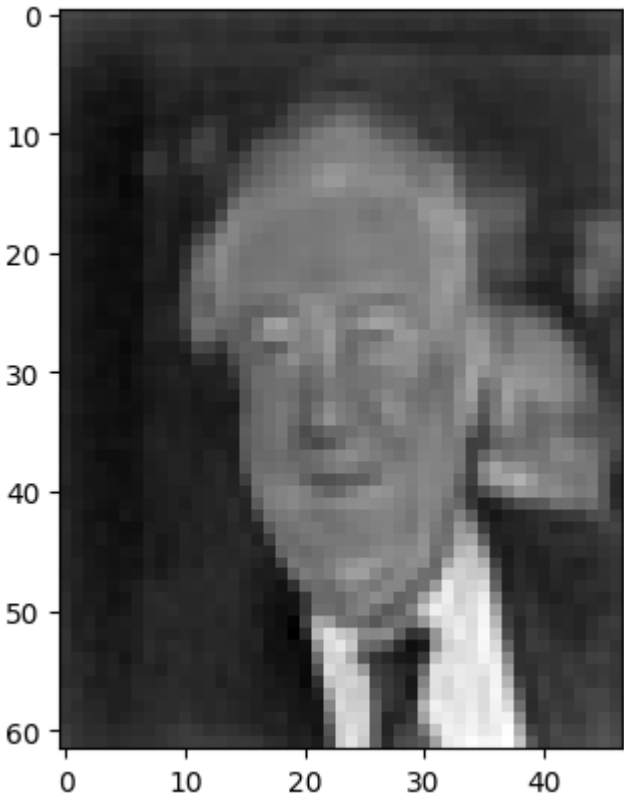
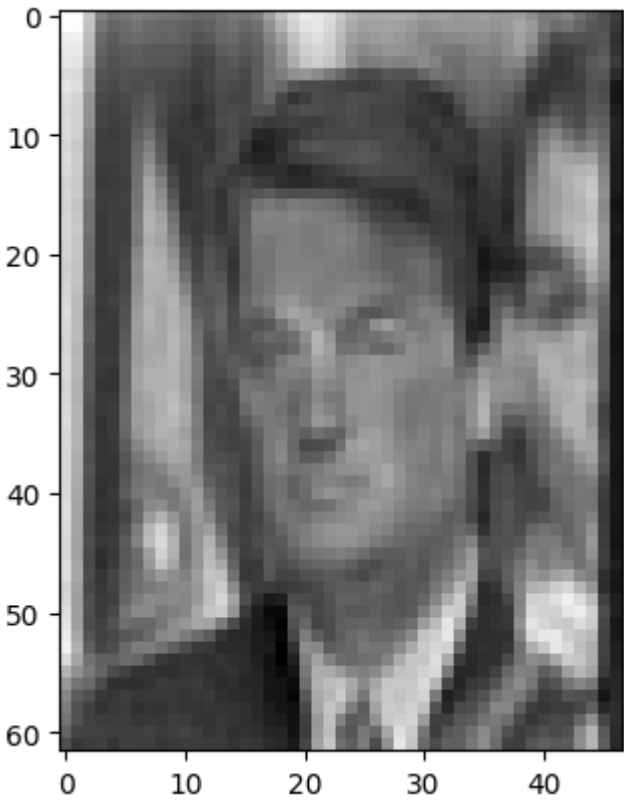


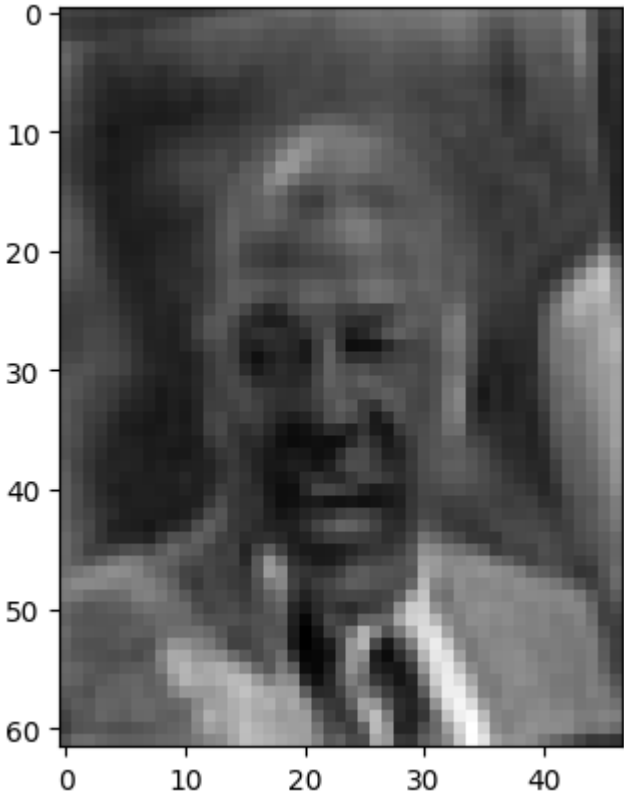



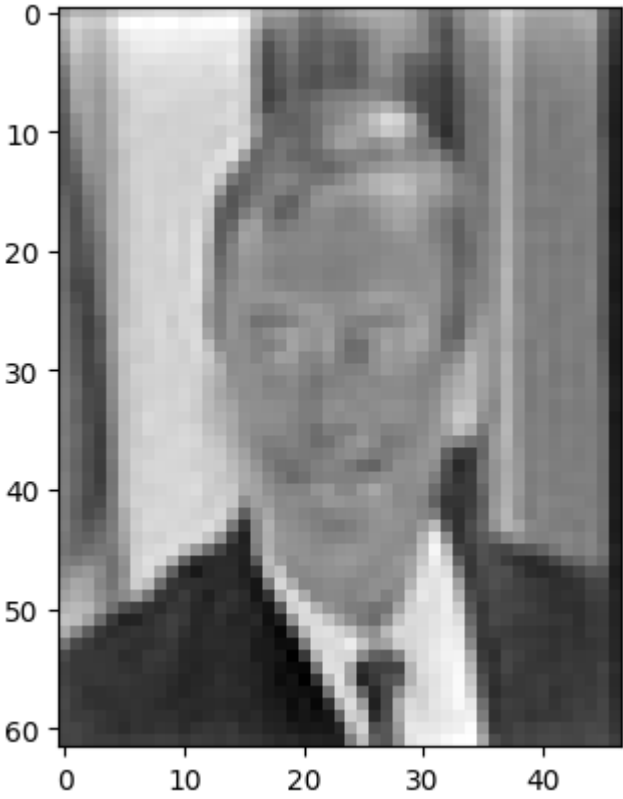






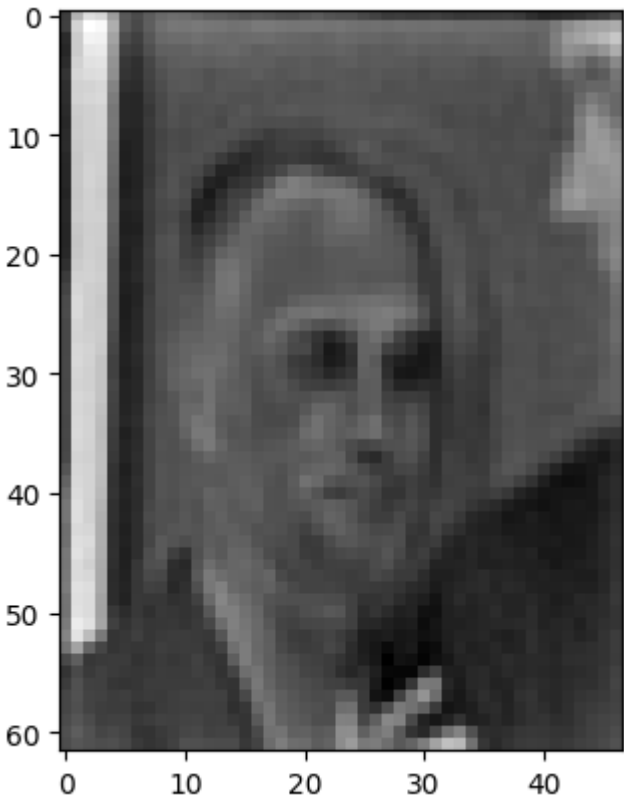


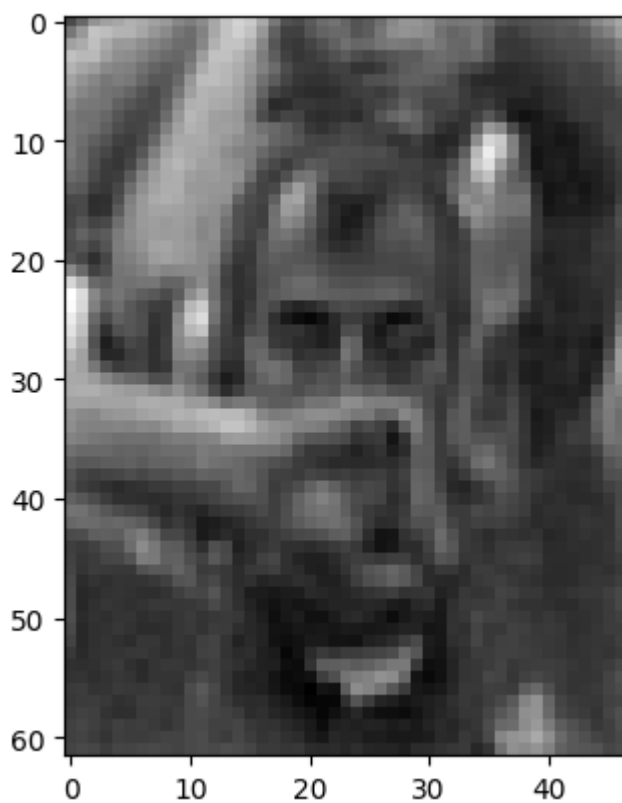













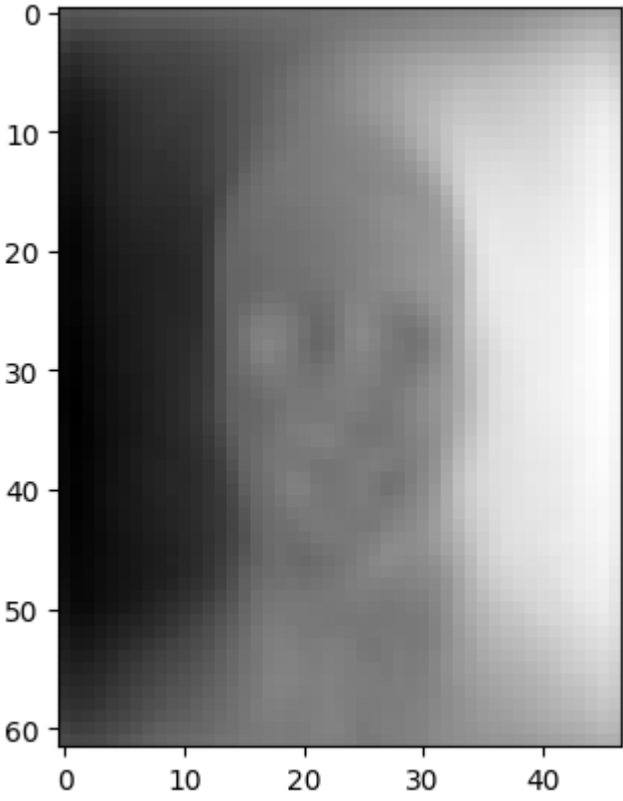
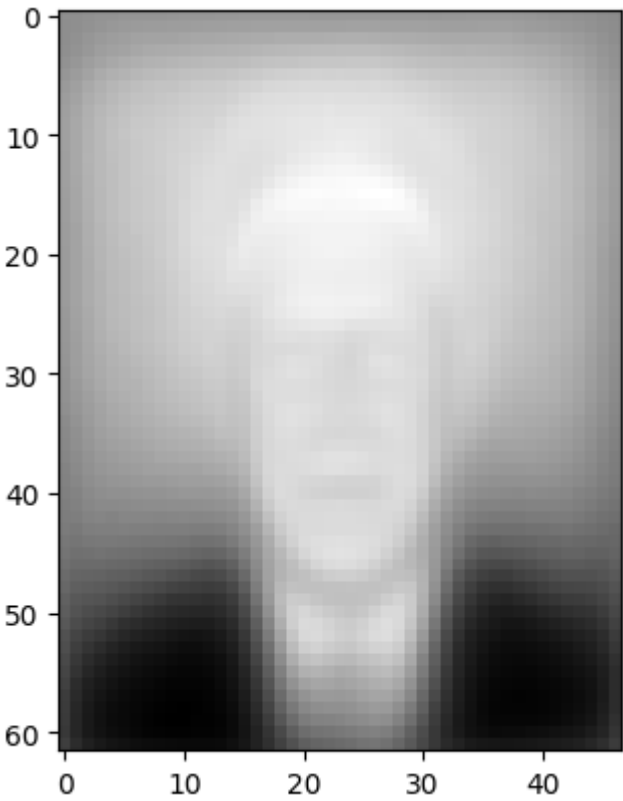


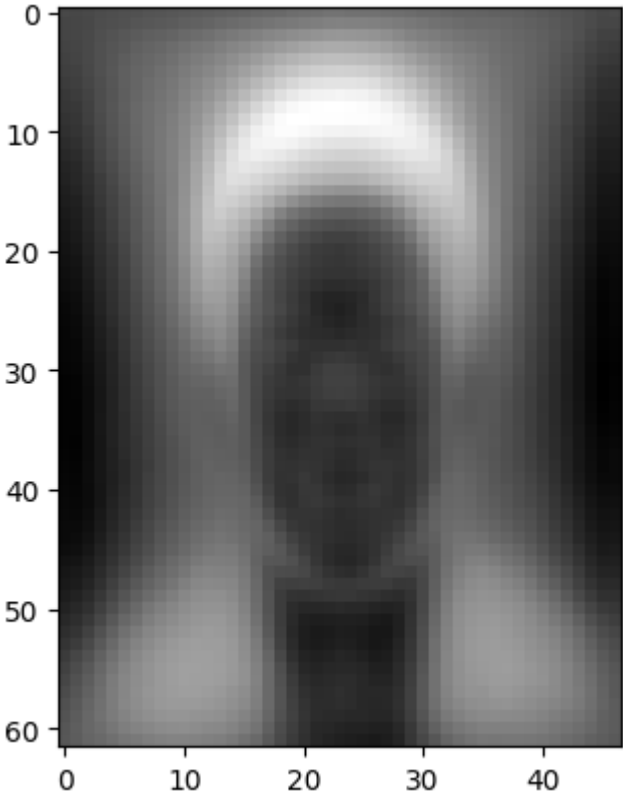

```
In [ ]: display(Latex(r"\newpage"))
```

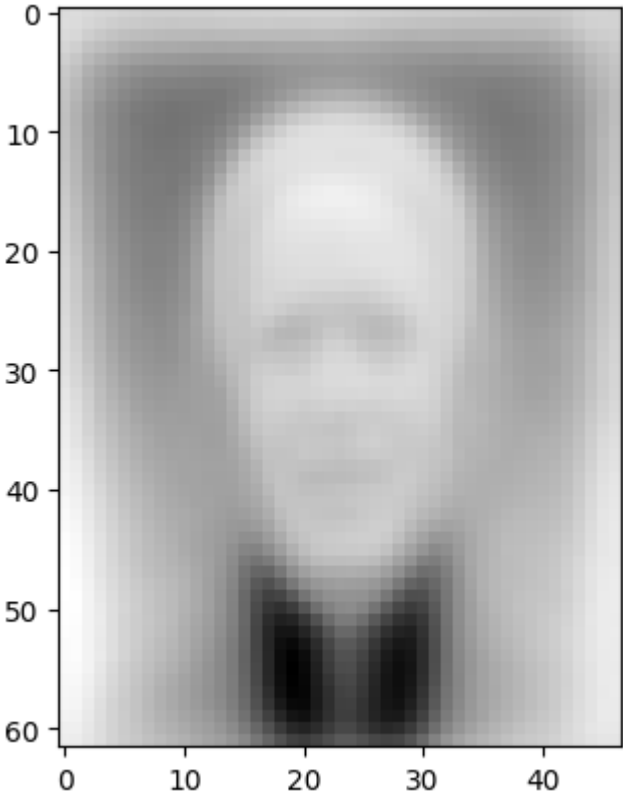

\newpage

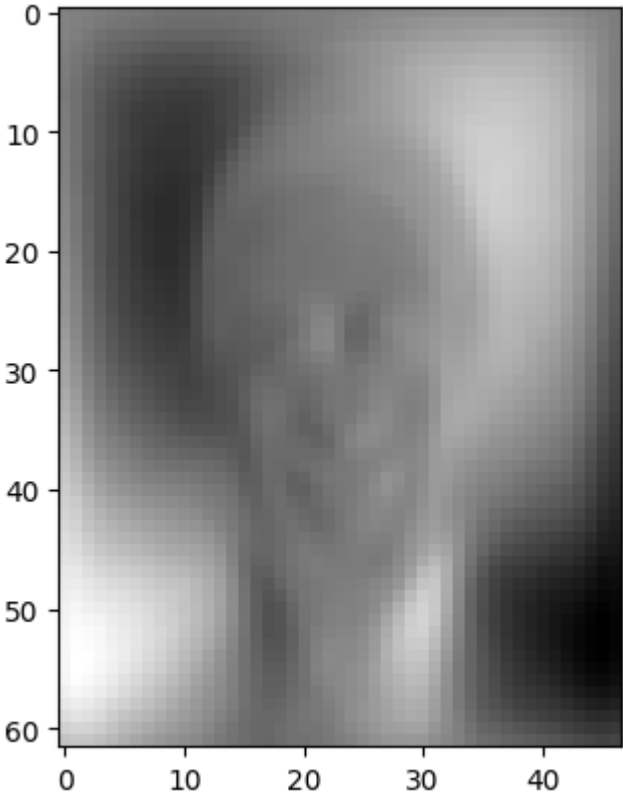
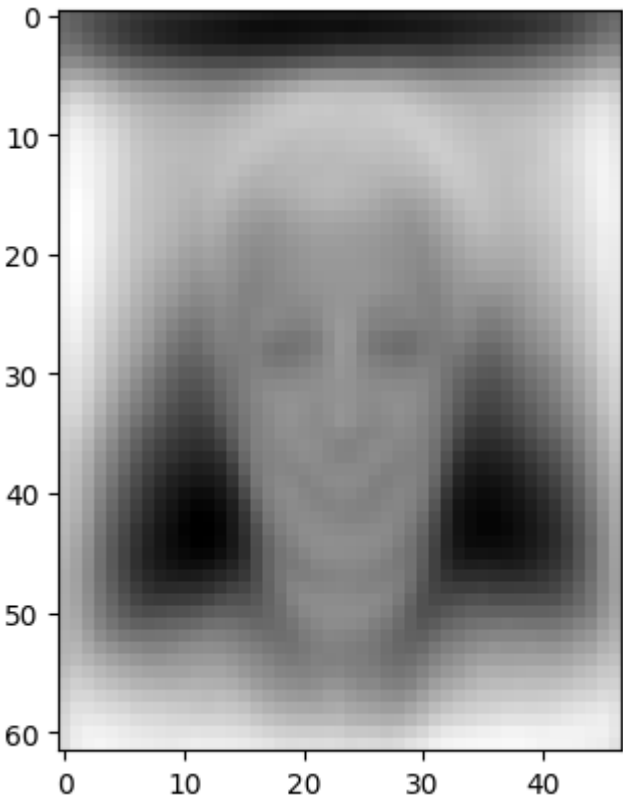
```
In [ ]: #1.d
        for i in range(20):
            matplotlib.pyplot.imshow(eig_vectors[:,i].reshape(62,47), cmap=matplotlib.cm.gray)
            matplotlib.pyplot.savefig(f'files/eigenface_nr.{i+1}.png')
            matplotlib.pyplot.show()
```







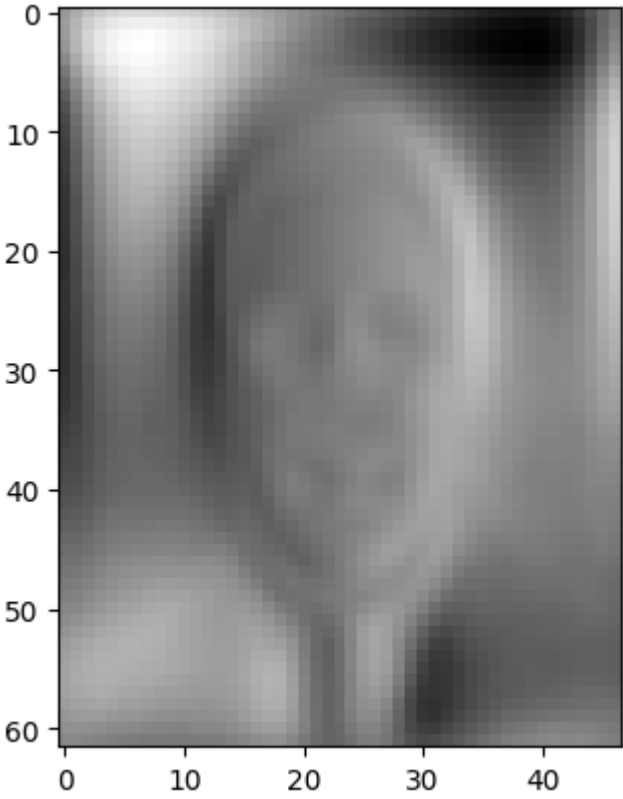







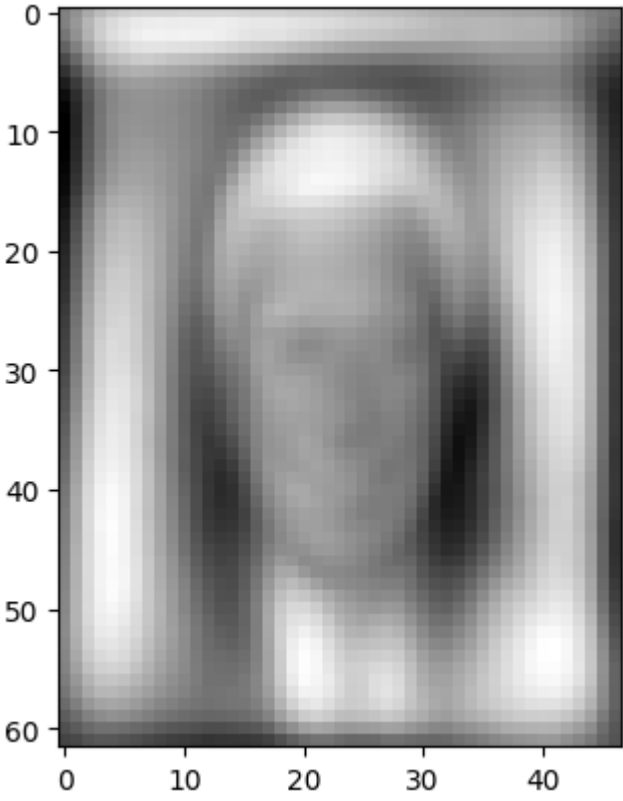
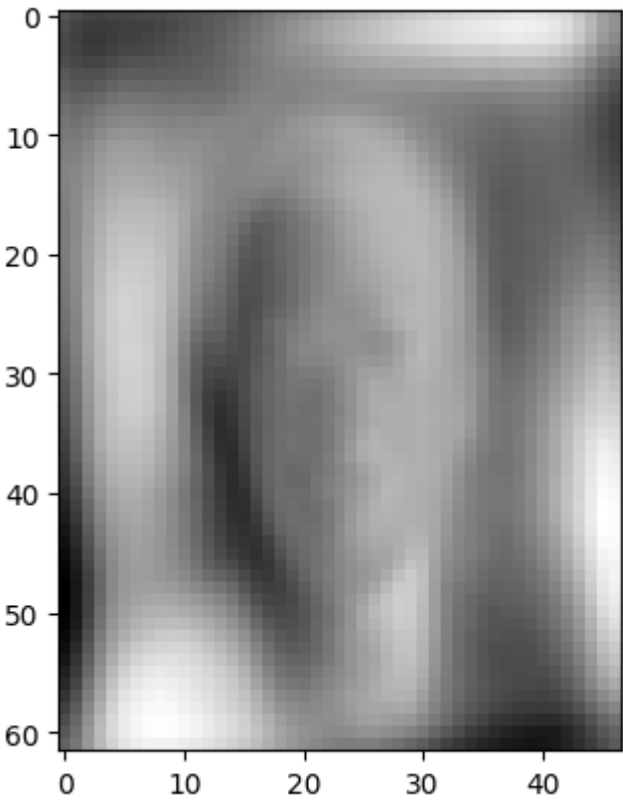


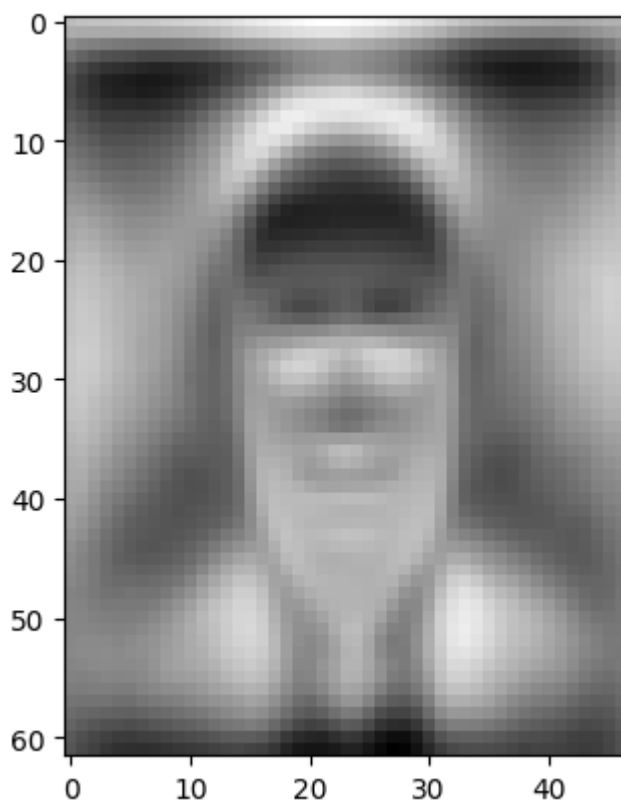













```
In [ ]: display(Latex(r"\newpage"))
```

\newpage

```
In [ ]: #1.e
```

```
explained_variances = []  
for i in range(len(eig_values)):  
    explained_variances.append(eig_values[i] / np.sum(eig_values)*100)  
variance_explained_cum = np.cumsum(explained_variances)
```

```


k = 0
PCDs_needed=[]
for i in range(len(explained_variances)):
    k+=explained_variances[i]
    PCDs_needed.append(explained_variances[i])
    if k>95:
        break

print(f"One can see that we need {len(PCDs_needed)} PCDs to explain 95% of the vari
matplotlib.pyplot.plot(np.arange(1, X.shape[1] + 1, 1), variance_explained_cum)
matplotlib.pyplot.ylabel('Variance Explained in percent')
matplotlib.pyplot.xlabel('Number of PCDs')
matplotlib.pyplot.title('Variance explained by number of PCDs')

```

One can see that we need 75 PCDs to explain 95% of the variance.

Out[ ]: Text(0.5, 1.0, 'Variance explained by number of PCDs')



In [ ]: `display(Latex(r"\newpage"))`

\newpage

```

In [ ]: #1.f
#load X again since our previous X was centered
X = dataset['data']

X_train = X[:int(X.shape[0]*0.8)]
X_test = X[int(X.shape[0]*0.8):]

eig_vectors, eig_values, _ = np.linalg.svd(X_train.T@X_train)

idx = eig_values.argsort()[::-1]
eig_values = eig_values[idx]
eig_vectors = eig_vectors[:, idx]

```

```

training_loss = []
test_loss = []

number_of_PCDs = [10, 20, 50, 100, 500, 1000, 2914]

for k in number_of_PCDs:
    k_eig_vectors = np.zeros(eig_vectors.shape)
    k_eig_vectors[:, :k] = eig_vectors[:, :k]

    X_train_reconstruction = X_train @ k_eig_vectors @ k_eig_vectors.T
    training_loss.append(1/(X_train.shape[0]*X_train.shape[1])*np.linalg.norm(X_train_reconstruction - X_train, 'fro'))

    X_test_reconstruction = X_test @ k_eig_vectors @ k_eig_vectors.T
    test_loss.append(1/(X_test.shape[0]*X_test.shape[1])*np.linalg.norm(X_test_reconstruction - X_test, 'fro'))

    print(f"For {k} number of PCDs, the training loss was: {training_loss[-1]}, and the test loss was: {test_loss[-1]}")

```

For 10 number of PCDs, the training loss was: 0.024876154512220056, and the test loss was 0.025062690624474873

For 20 number of PCDs, the training loss was: 0.01832719678080922, and the test loss was 0.018420726054348077

For 50 number of PCDs, the training loss was: 0.01145844974748309, and the test loss was 0.011592239534264808

For 100 number of PCDs, the training loss was: 0.00749010035300871, and the test loss was 0.007688371896088376

For 500 number of PCDs, the training loss was: 0.0016752821280762095, and the test loss was 0.001950320561979342


For 1000 number of PCDs, the training loss was: 0.0005270785648694966, and the test loss was 0.0007230311720999202

For 2914 number of PCDs, the training loss was: 2.313802221591989e-16, and the test loss was 2.8704234588938264e-16

```

In [ ]: #1.f
#the plot
matplotlib.pyplot.plot(number_of_PCDs, training_loss, label = "Training loss")
matplotlib.pyplot.plot(number_of_PCDs, test_loss, label = "Test loss")
matplotlib.pyplot.title("Training and test loss for different number of PCDs")
matplotlib.pyplot.xlabel("Number of PCDs")
matplotlib.pyplot.ylabel("Loss")
matplotlib.pyplot.legend()
matplotlib.pyplot.show()

```



I worked with:

- Jens Kristoffersen
- Dherick Derahmanan
- Olan Nomeland
- Samhet Behara

I certify that all solutions are entirely in my own words and that I have not looked at another student's solutions. I have given credit to all external sources I consulted.



2. a

Building the heap:  $O(k)$  where  $k$  is the size of the heap.

Remaining  $n-k$  points take  $O(\log k)$  time.

Finding the majority label of the  $k$  closest neighbours takes  $O(k)$  time.

Finding  $n$  euclidian distances takes  $n O(d)$  time.

$$\text{Total: } O(k) + (n-k)O(\log k) + O(k) + n O(d) = O(nd) + O(k) + O(n \log k)$$

2.6

Using the problem of placing balls in boxes, the number of monomials is  $\binom{d+p}{p}$  and this is the dimension the point lives in.

$$d \text{ is now } \frac{(d+p)!}{(d+p-p)!p!} = \frac{(d+p)!}{d!p!}$$

New runtime:

$$O\left(n \cdot \frac{(d+p)!}{d!p!}\right) + O(k) + O(n \log k)$$

2.C

1D: left and right = 2

2D: all sides and all corners  
= 8

dD: There are  $2 \cdot d$  adjacent  
cells and  $2^d$  corners  
=  $2d + 2^d$

You would need to  
check each one of these cells so  
time complexity is  $O(2d + 2^d)$

3.a

I will be 0 surprised since  
I knew with 100% certainty  
what the outcome would be.

3.6.

Max surprised, my prior belief was that there was a 0% chance of picking up a white ball.

3.C

When  $P_B$  is either 0 or 1  
is when the entropy is  
minimized.

$$H_b(0) = -0 \log 0 - 1 \log 1 \\ = \underline{\underline{0}}$$


Maximized at  $P_B = 0.5$

$$H_b(0.5) = -0.5 \log 0.5 - 0.5 \log 0.5 \\ = \underline{\underline{1}}$$

3. d

↑  
1

entropy



5 hrs Ltg

con can e

3. e.

$$\begin{aligned} P &= P(Y=1 | X_{j,v}=1) P(X_{j,v}=1) \\ &\quad + P(Y=1 | X_{j,v}=0) P(X_{j,v}=0) \\ &= q_1 P(X_{j,v}=1) + q_2 P(X_{j,v}=0) \\ &= \lambda q_1 + (1-\lambda) q_2 \text{ where} \\ &\quad \lambda \text{ is the probability that} \\ &\quad X_j \leq v. \end{aligned}$$

show that  $H(Y) - H(Y | X_{j,v})$   
is always positive if  $q_1 \neq q_2$

$$\begin{aligned} \rightarrow H_b(P(Y=1)) - \sum P(X_{j,v}=i) H_b(P(Y=1 | X_{j,v}=i)) \\ > 0 \end{aligned}$$

$$\rightarrow H_b(\lambda q_1 + (1-\lambda) q_2) > \lambda \cdot H_b(q_1) + (1-\lambda) H_b(q_2)$$

$$\begin{aligned} \rightarrow -(\lambda q_1 + (1-\lambda) q_2) \log(\lambda q_1 + (1-\lambda) q_2) \\ - (\lambda q_1 + (1-\lambda) q_2) \log(\lambda q_1 + (1-\lambda) q_2) > \end{aligned}$$



$$\lambda \cdot (-q_1 \log q_1 - (1-q_1) \log(1-q_1)) \\ + (1-\lambda) \cdot (-q_2 \log q_2 - (1-q_2) \log(1-q_2))$$

→ sees that as  $H_b$  is strictly concave, multiplying with lambda on the outside of  $H_b$  function will lead to a smaller value.

I do not have a better explanation i

3. f

Average:

$$\begin{aligned} E\left[\frac{1}{n} \sum_{i=1}^n Y_i\right] &= \frac{1}{n} \sum_{i=1}^n E[Y_i] \\ &= \frac{1}{n} \cdot \sum_{i=1}^n \mu = \frac{1}{n} \cdot n\mu = \mu \end{aligned}$$

Variance:

$$\begin{aligned} \text{Var}\left(\frac{1}{n} \sum_{i=1}^n Y_i\right) &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(Y_i) \\ &= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n} \end{aligned}$$

3.9.1

Prob of not being picked:

$$\left(\frac{n-1}{n}\right)^n.$$

Lowest value of this function is at  $n=25$  since it is monotonically increasing.

At  $n=25$ :

$$\left(\frac{25-1}{25}\right)^{25} = 0.3604 \approx 0.36$$

And at its limit:

$$\lim_{n \rightarrow \infty} \left(\frac{n-1}{n}\right)^n = \lim_{n \rightarrow \infty} \left(\frac{n(1-\frac{1}{n})}{n}\right)^n$$

$$= \lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n = \lim_{n \rightarrow \infty} e^{n \ln(1-\frac{1}{n})}$$

$$= \frac{1}{e} \quad \square \rightarrow P \text{ lies in } [0.36, \frac{1}{e}]$$

3.g.ii

More trees  $\rightarrow$  more computation

Plot the loss at different  
nr. of trees in an increasing  
manner. and pick the number  
of trees when the loss levels  
sufficiently out towards its  
asymptote.

3.6

$$\text{Var}\left(\frac{1}{n} \sum_{i=1}^n Z_i\right) = \frac{1}{n^2} \text{Var}\left[\sum_{i=1}^n Z_i\right]$$

$$= \frac{1}{n^2} \left( \sum_{i=1}^n \text{Var}(Z_i) + 2 \sum_{i < j}^n \text{Cov}(Z_i, Z_j) \right)$$

$$= \frac{1}{n^2} \left( n\sigma^2 + 2\rho \frac{(n-1)n}{2} \right)$$

$$= \frac{\sigma^2}{n} + \frac{\rho(n-1)}{n} \quad \square$$

4

```
In [ ]: from IPython.display import Latex
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]: #Import dataset
import sklearn
from sklearn import datasets
diabetes = sklearn.datasets.load_diabetes()
```

```
In [ ]: display(Latex(r"\newpage"))
```


\newpage

```
In [ ]: #4.a
from sklearn.manifold import TSNE
x = diabetes["data"]
y = diabetes["target"]
d = 2
tsne = TSNE(d)
tsne_result = tsne.fit_transform(x)

points = plt.scatter(tsne_result[:, 0], tsne_result[:,1], c = y, cmap = 'viridis')
plt.colorbar(points)
```

c:\Users\elias\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\manifold\\_t\_sne.py:800: FutureWarning: The default initialization in TSNE will change from 'random' to 'pca' in 1.2.  
 warnings.warn(  
c:\Users\elias\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\manifold\\_t\_sne.py:810: FutureWarning: The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.  
 warnings.warn(

```
Out[ ]: <matplotlib.colorbar.Colorbar at 0x1cff409fca0>
```




```
In [ ]: display(Latex(r"\newpage"))
```

\newpage

4.b The feature is sex.

```
In [ ]: points = plt.scatter(tsne_result[:, 0], tsne_result[:,1], c = x[:,1], cmap = 'viridis')
plt.colorbar(points)
```

```
Out[ ]: <matplotlib.colorbar.Colorbar at 0x1cff3ed0460>
```



```
In [ ]: display(Latex(r"\newpage"))
```

\newpage


```
In [ ]: #4.c
new_x = np.delete(diabetes["data"], axis = 1, obj = 1)
tsne = TSNE(d)

new_result = tsne.fit_transform(new_x)

points = plt.scatter(new_result[:, 0], new_result[:,1], c = y, cmap = 'viridis')
plt.colorbar(points)
```

```
c:\Users\elias\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\manifold\_t_sne.py:800: FutureWarning: The default initialization in TSNE will change from 'random' to 'pca' in 1.2.
  warnings.warn(
c:\Users\elias\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\manifold\_t_sne.py:810: FutureWarning: The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.
  warnings.warn(
```

```
Out[ ]: <matplotlib.colorbar.Colorbar at 0x1cff3f6ffd0>
```



We do not see any two clear clusters anymore since tsne does not differentiate on sex. In other words, most of the data points got a lot of their variance from sex meaning females and males are quite similar besides being of different genders. The fact that sex was included ended up with "pulling" the data apart and forming two clusters.

If the change did not occur and we did still see two clear clusters, it would mean that there would be another feature which was very binary in its distribution. If age was split in above 50 or below 50 this would also form two clusters.


```
In [ ]: display(Latex(r"\newpage"))
```



\newpage

```
In [ ]: #4.d
from sklearn.decomposition import PCA
pca = PCA(d)
pca_result = pca.fit_transform(x)
points = plt.scatter(pca_result[:, 0], pca_result[:,1], c = y, cmap = "viridis")
plt.colorbar(points)
```

Out[ ]: <matplotlib.colorbar.Colorbar at 0x1cff42b4220>



```
In [ ]: display(Latex(r"\newpage"))
```

\newpage

```
In [ ]: #4.e
def MSE(y_hat, y_true):
    return (1/y_hat.shape[0])*np.sum((y_hat - y_true)**2)

X_train = x[:100]
y_train = y[:100]
X_test = x[100:]
y_test = y[100:]

X_train_mtx = np.hstack([X_train, np.ones((100,1))])
X_test_mtx = np.hstack([X_test, np.ones((342, 1))])
w_ols = np.linalg.inv(X_train_mtx.T @ X_train_mtx) @ X_train_mtx.T @ y_train
y_hat = X_test_mtx@w_ols

test_mse = MSE(y_hat, y_test)
print(f"The test MSE is: {test_mse}")

def c_index(y_hat, y_test):
    nr_conc = 0
    nr_disc = 0
```

```

for i in range(y_hat.shape[0]):
    for j in range(y_test.shape[0]):
        if i == j:
            continue
        else:
            y_test_i = y_test[i]
            y_test_j = y_test[j]
            y_hat_i = y_hat[i]
            y_hat_j = y_hat[j]
            if y_test_i > y_test_j and y_hat_i > y_hat_j:
                nr_conc += 1
            elif y_test_i > y_test_j and y_hat_i < y_hat_j:
                nr_disc += 1
    return nr_conc/(nr_conc+nr_disc)

print(f"The c index is: {c_index(y_hat, y_test)}")

```

The test MSE is: 3430.9233826005243

The c index is: 0.7452930850514576

In [ ]: `display(Latex(r"\newpage"))`

\newpage

```

In [ ]: #4.f
from sklearn.linear_model import Ridge
def cross_validation(X, y, k, model, loss_calculator):
    n = X.shape[0]
    idx = np.random.permutation(n)
    X_shuffled = X[idx]
    y_shuffled = y[idx]

    size = n // k

    validation_loss = np.zeros((n,))
    training_loss = np.zeros((n,))

    for i in range(k):
        start = i*size
        end = (i+1)*size

        X_test = X_shuffled[start:end]
        y_test = y_shuffled[start:end]

        X_train = np.vstack([X_shuffled[:start], X_shuffled[end:]]
        y_train = np.hstack([y_shuffled[:start], y_shuffled[end:]]

        model.fit(X_train, y_train)
        training_loss[i] = loss_calculator(model.predict(X_train), y_train)
        validation_loss[i] = loss_calculator(model.predict(X_test), y_test)

    return np.mean(validation_loss), np.mean(training_loss)

lambdas = [10**-5, 10**-4, 10**-3, 10**-2, 10**-1, 1]
validation_MSE = []
training_MSE = []

for l in lambdas:

```

```

validation_MSE_current, training_MSE_current = cross_validation(X_train, y_train,
validation_MSE.append(validation_MSE_current)
training_MSE.append(training_MSE_current)

print(f"Lambda={l} has validation loss={validation_MSE_current} and training loss={training_MSE_current}")
print("")

plt.plot(np.log10(lambdas), training_MSE, label='Training MSE')
plt.plot(np.log10(lambdas), validation_MSE, label='Validation MSE')
plt.legend()
plt.xlabel('log(Lambda)')
plt.ylabel('MSE')

best_lambda = lambdas[np.argmin(validation_MSE)]
ridge_best_lambda = Ridge(alpha=best_lambda).fit(X_train, y_train)
ridge_y_hat = ridge_best_lambda.predict(X_test)
ridge_mse = MSE(ridge_y_hat, y_test)
print(f"The best lambda was {best_lambda}.\nThis lambda gave MSE with ridge = {ridge_mse}")

```

Lambda=1e-05 has validation loss=404.15799243503966 and training loss=385.8010227690261

Lambda=0.0001 has validation loss=416.3534904067575 and training loss=385.06213726177435

Lambda=0.001 has validation loss=403.72113028383114 and training loss=385.94238469552073


Lambda=0.01 has validation loss=402.9057175061709 and training loss=385.8414729654401

Lambda=0.1 has validation loss=402.7489796064022 and training loss=385.87803508971024

Lambda=1 has validation loss=401.3585709192019 and training loss=385.9097279767526

The best lambda was 1.

This lambda gave MSE with ridge = 5039.062537574326.



In [ ]: